

# C++

April 11, 2022

Mike Spertus

[spertus@uchicago.edu](mailto:spertus@uchicago.edu)



MASTERS PROGRAM IN  
COMPUTER SCIENCE

THE UNIVERSITY OF CHICAGO



# Projects

- Consider starting to think about a project choice
  - No need to start coding yet, but you are encouraged to come up with an idea in the next 1-2 weeks
  - Feel free to consult with me if you aren't sure what will be suitable
- No requirements other than demonstrating your facility with a number of the techniques taught during this course
  - It doesn't need to cover all of them
    - E.g., if you have a lot of advanced template code that could compensate for not having advanced concurrency



# ADVANCED VARIADICS



# Variadics

- A variadic template is a template that can take a variable number of parameters called a parameter pack
- Here's how to create a function that adds up all of its arguments

```
template<typename T>
T adder(T v) {
    return v;
}
```

```
template<typename T, typename... Args>
std::common_type_t<T, Args...> adder(T first, Args... args)
{
    return first + adder(args...);
}
```

- For example, `adder(1, 2, 3, 4)` will be 10, and `adder("foo"s, "bar"s)` returns the string "foobar"s
  - Reminder: The `s` after the closing double-quote is a string literal



# Pack Expansion

- ... is used to expand parameter packs into comma-separated lists (roughly)
- In the previous slide, we expanded args... to 2, 3, 4



# What can be expanded by ...?

- ... doesn't have to be applied directly to a pack
- It can be applied to any expression that contains a pack
- $(\text{args} + 2)\dots$  expands to 3, 4, 5
- <https://godbolt.org/z/dfdfjfKqW>

# Expanding an expression with variadic parameter packs



- Everything inside the expression gets expanded and separated by commas
- ```
// A version that takes its arguments by const
reference
template<typename T>
T adder_ref(T const &v) {
    return v;
}
// Args const &... expands each argument,
// puts a const & after it and separates
// it with commas
template<typename T, typename... Args>
common_type_t<T, Args...>
adder_ref(T const &first, Args const &... args) {
    return first + adder(args...);
}
```

# What if there are two packs in an expansion



- If multiple packs appear in an expansion, they are expanded in "parallel"
- They must have the same length
- $(args^*args)...$  expands to 1, 4, 9





# Wait, we've seen that before!

- This "simultaneously expanding internal packs in expressions" is what makes perfect forwarding work!
- Consider the implementation of `make_unique` from last quarter
- ```
template<typename T, typename ...Args>
unique_ptr<T> make_unique(Args&& ...args) {
    return new T(forward<Args>(args)...);
}
```

# Variadics case study: Improved printing



- Nearly every program does formatted output
- Yet the two built-in ways provided by C++, `IOStreams` and `printf` have serious problems
- We will review the problems described in the prelecture
- and then show how to fix them with variadics



# The problem with IOStreams

- C++ replaced “printf”-based I/O with I/O streams like you learned
- The problem is that IOStreams are very cumbersome for many simple tasks
- For example, we often see code like:
  - ```
} catch(MyException e) {  
    ostringstream msg;  
    msg << "Operation "  
        << e.operation  
        << " failed with error code "  
        << e.errorCode;  
    Logger.output(msg.str()); // Some logging framework  
}
```
  - In fact, what you more often see in this case, is a lame “Error detected” message to avoid the 6 lines above
  - We'd rather write  

```
printf("Operation %s failed with error code %d", e.operation, e.error_code);
```



# The problems with printf

- Unsurprisingly, many people continue to use printf with C++
- However, printf is not extensible
  - It doesn't know how to print any new types you have created
- Even worse, it ignores argument types and prototypes and crashes and corrupts the stack at runtime

# This code compiles without warning but crashes



- ```
int main()
{
    printf("%s\n", 1, 2);
}
```
- The problem is that unlike nearly every other function, `printf` doesn't declare or check the type or number of arguments it was called with because that depends on the format string
- ```
int printf(char *fmt, ...);
```
- The `...` could be anything
  - Don't confuse with C++ variadics



# Solution: C++20 format

- We can use variadic templates to create a typesafe extensible formatting library that is easier to use than printf
- `cout << format("id {} - name {}", 5, "Mike");`
- Automatically uses the right formatter
  - So I don't need to specify the type in the format string like I did with printf



# Implement with variadics

- Let's practice with variadics by implementing a simplified version of format ourselves
- <https://godbolt.org/z/bnn7ss33s>



# First handle the base case

- A typical approach in variadics is to have a non-variadic "base case" overload and then recurse
- This case is easy for format
  - No arguments to format!
- ```
string simple_format(string_view fmt) {  
    return string(fmt);  
}
```





# How to print an argument

- As we saw in the prelecture, formatters for format are really complicated
  - And not really relevant for this variadics discussion
- We'll just use the stream inserter to print a type
- ```
template<typename T>
string make_string(T const &t) {
    ostringstream oss;
    oss << t;
    return oss.str();
}
```
- IRL this would be too slow and clunky, but...



# Now recurse and we're done

```
template<typename T, typename ...Ts>
string simple_format(string_view fmt, T const&t, Ts const &... ts)
{
    auto braces = fmt.find("{}"); // Omitting error handling :(
    return string(fmt.substr(0, braces)) + make_string(t)
        + simple_format(fmt.substr(braces + 2), ts...);
}
```



# Fold operations

- This kind of recursing down a variadic parameter pack is very common
- But it can be verbose and clunky for simple things
- Look at our adder implementation
- C++17 added fold operations to simplify
- You can think of fold operations as a variadic analog to accumulate



# Fold operations

- Putting the `...` before a binary operator followed by something containing a pack puts the operator between every element in the pack
- Sounds complicated but really isn't
- For example, if `args...` is 1, 2, 3, then  
    `... + args`  
evaluates to  
    `1+2+3`



# adder redux

- Now, we can write adder simply without many levels of recursion
- ```
template<typename... Args>
common_type_t<Args...>
adder(Args... args) {
    return (... + args);
}
```

# There are some variations on fold expressions



- In addition to the version shown here
- There are variants of fold expressions that
  - Start with an initial "accumulator" value
  - Fold from the right rather than the left
- They otherwise work the same as what we did and should be straightforward to pick up if you need them



# Tuples

- We've been using tuples a lot
- Now we have the tools to understand them more deeply
- Tuples are a generalization of `std::pair` to any number of fields and the most important and prevalent example of variadics
- ```
tuple<string, int, double> si("str", 2, 3.5);  
cout << get<0>(si) // prints "str"  
cout << get<int>(si); // prints 2 (C++14)  
int two = get<1>(si);
```

# Why isn't get a member function?



- In the previous slide, why do we say `get<1>(di)` instead of the more natural `di.get<1>()`?
- Consider the following:

```
template<typename ...T>
auto f(tuple<T...> t)
{ return t.get<0>(); }
```
- The compiler might think that `t.get` is a field in `t` and interpret it as asking whether `t.get` is less than `0`!
- In order to disambiguate, you would have to do the ugly

```
template<typename ...T>
auto f(tuple<T...> t)
{ return t.template get<0>(); }
```
- Kind of reminiscent of the ugly `typename` disambiguator mentioned above
- Since we didn't want such an important basic vocabulary class to require using such an advanced technique, we made `get` into a global function



# Let's try to implement tuples



- Tuple.h



# Returning multiple values

- We already learned about using tuples to return multiple values at runtime
- `pair<int, int> f() { return {1, 2}; // ok }`
- `tuple<int, int, char> f() { return {1, 2, 'u'}; } // OK in C++17`
- `auto [i, j, c] = f(); // Structured binding back into variables`

# A purely compile-time use of tuples



- Often, one uses tuples as a sort of compile-time “typelist”
  - These tuples are not necessarily meant to be instantiated at run-time
  - Instead, we are using the tuple type as a convenience container of types at compile-time
- The next few slides will show how to manipulate typelists
- We will use these techniques in our tuple implementation
- For illustration, assume that the following types have been defined.
  - Note that we use tuple just a “compile-time container of types” to make it easier to organize variadic template arguments

```
struct A {  
    static void foo() { cout << "In A" << endl; }  
};  
struct B {  
    static void foo() { cout << "In B" << endl; }  
};  
using AB = tuple<A, B>;  
Using CharTypes = tuple<char, unsigned char, signed char>;
```



# Calculating length

```
template<class TList> struct Length;
template<>
struct Length<tuple<>>
{
    static size_t constexpr value = 0;
};
template<class T, typename... Us>
struct Length<tuple<T, Us...> >
{
    // Of course, we could just use sizeof...(Us),
    // but let's see how to do ourselves
    static size_t constexpr value
        = 1 + Length<tuple<Us...>>::value };
};
int main() // Prints 3
{
    cout << Length<CharTypes>::value << endl;
}
```

# Finding the type at a given index



```
template<class List, int i> struct TypeAt;
template<class Head, typename... Tail>
struct TypeAt<tuple<Head, Tail...>, 0>
{
    using type = Head;
};
template<class Head, typename... Tail, int i>
struct TypeAt<tuple<Head, Tail...>, i>
    : public TypeAt<tuple<Tail...>, i - 1> // Inherits this result
{
};

int main() // prints "In B"
{
    // Note that I have to say TypeAt<xxx>::type to get the result
    TypeAt<AB,1>::type::foo();
}
```

# Finding the index of a given type



```
template<class List, class Target> struct IndexOf;
template<class Target>
struct IndexOf<tuple<>, Target>
{
    static size_t constexpr value = -1; // Return -1 if not found
};
template<class ...Tail, typename Target>
struct IndexOf<tuple<Target, Tail...>, Target>
{
    static size_t constexpr value = 0;
};
template<class Head, typename... Tail, class Target>
struct IndexOf<tuple<Head, Tail...>, Target>
{
private: // Using a Compile-time temporary
    static size_t constexpr temp = IndexOf<tuple<Tail...>, Target>::value;
public:
    static size_t constexpr value = temp == -1 ? -1 : 1 + temp;
};

int main() // Prints 1
{
    cout << IndexOf<CharTypes, unsigned char>::value;
}
```

# Appending the types in two tuples



```
template<class First, class Second> struct Append;
```

```
template<typename... Ts, typename... Us>
struct Append<tuple<Ts...>, tuple<Us...> >
{
    using type = tuple<Ts..., Us...>;
};
```

```
int main() // Prints 2
{
    typedef Append<AB, tuple<int>>::type ABInt;
    cout << IndexOf<ABInt, int>::value;
}
```



# Replacing a type in a tuple

```
template<typename T, typename A, typename B> struct Replace;
```

```
template<typename... Ts, typename A, typename B>  
struct replace<tuple<A, Ts...>, A, B> {  
    using type = tuple<B, Ts...>;  
};
```

```
template<typename H, typename... Ts, typename A, typename B>  
struct Replace<tuple<H, Ts...>, A, B>  
    : public Append<tuple<H>, typename Replace<tuple<Ts...>, A, B>::type> {  
};
```

- Why did I have to say `typename` in front of `Replace<Replace<tuple<Ts...>, A, B>::type>`?
- The point is that when the compiler is looking at the above specialization, it doesn't know what types will be in `Ts...`, so it can't assume that `Replace<tuple<Ts...>, A, B>::type` is a type or not. (For example, it might be a static member variable). We clarify this to the compiler by preceding it with the keyword `typename`.
- If you are getting compiler errors because the compiler can't figure out that something is a type, just preceded it with the term "`typename`"





# The Empty Base Optimization

- In C++, the address of an object is also the identity of the object
- For example, if `a1` and `a2` are two different objects, then they should have different addresses
- This could break if they have size 0 (e.g., an array of size 0 objects)
- So even an empty class has size  $> 0$ !
- However, an empty base class does not take space since the total object addresses
  - Food for thought: Does it always address?

# Simulating template virtuals



- We will put together a lot of template techniques for a really mind-bending payoff
- We will need one more tool at our disposal
- Just like regular methods, we sometimes wish that method templates were virtual
- However, method templates cannot be virtual (Why?)
- Let's see if we can come up with a way to simulate virtual method templates

# What would template virtuals look like?



- The reason method templates can't be virtual is because they aren't methods
  - What would the vtable even look like?
- However, we can call ordinary methods from our method templates
- So what we will do is specify which specializations should call which method

# What interface do we want to support



- Suppose we have a method template
- ```
struct S {  
    template<typename T> void f();  
};
```
- Now suppose we want `f<int>` and `f<double>`, and `f<A>` to be “virtuals” that there is a way to override them in a derived class (A is just another class)



# First swing

- A first attempt might be something like:

```
struct S {  
    template<class T>  
    void f() { return fHelper(T()); }  
    virtual void fHelper(int &&) = 0;  
    virtual void fHelper(double &&) = 0;  
    virtual void fHelper(A &&) = 0;  
};
```

- This actually works for the int and double case
- A class that inherits from S can override the fHelpers
- Unfortunately, A may be abstract or not have a default constructor, so fHelper(A ()) won't be able to create an A

# TT



- What we need is a “tag type” that can always be created
  - Not abstract
  - Default constructible
  - Cheap to construct
- ```
template<class T>
struct TT {
};
```
- Fits the bill. If I construct with `TT<A>{ }`, it’s legal because it is not abstract and is default constructible, and it’s cheap because it’s just an empty class



# Better

- Now, improve our previous approach:

```
struct S {  
    template<class T>  
    void f() { return fHelper(TT<T>()); }  
    virtual void fHelper(TT<int> &&) = 0;  
    virtual void fHelper(TT<double> &&) = 0;  
    virtual void fHelper(TT<A> &&) = 0;  
};
```

- A class that inherits from S can override the fHelper()s



# Can we automate?

- The only problem is that it is tedious and error prone to manually define all of the helper virtual methods
- Sounds like a job for variadics
- Let's begin by defining a class template that holds just one Helper
- And then use variadics to inherit from all of them





# The solution

```
template<typename T>
struct Holder {
    virtual void fHelper(TT<T> &&) =0;
};

template<typename ...Ts> // which holders to inherit
struct ST : public Holder<Ts>... {
    template<class T>
    void f() { return fHelper(TT<T>()); }
};

using S = ST<int, double, A>;
```

- S has virtual fHelpers for int, double, and A

# Wow, that seems obscure

## Why do I care?



- Well, you don't yet
- Since our mindbending application will need this
- For now, just suspend your disbelief



# HW 13-1

- Create a `ReplaceAll<TL, T, U>` template that replaces all occurrence of `T` in `TL` with `U`.
  - For example,  
`ReplaceAll<tuple<char, int, char>, char, double>::type`  
should be  
`tuple<double, int, double>`
- Define a `Reverse` template that reverses the parameters of a tuple
  - For example, `Reverse<tuple<A, B>>::type` will be `tuple<B, A>`



# HW 13-2

- Extend Tuple2 (on canvas) so that you can get elements by type as well as index
  - ```
Tuple2<int, double> t2id(4, 5.6);  
cout << "get<1>(t2id) = " << get<1>(t2id) << endl; // Works now  
cout << "get<int>(t2id) = " << get<int>(t2id) << endl; // HW
```



## HW 13-3: Extra Credit

- The "outer product" HW for last week took the outer product of vectors
- Why wouldn't the approach work for arrays?
  - E.g., `std::array<A, 5>`
- How would you make it work?
  - **Warning:** Surprisingly difficult