

C++

May 2, 2022

Mike Spertus

spertus@uchicago.edu



MASTERS PROGRAM IN
COMPUTER SCIENCE

THE UNIVERSITY OF CHICAGO



"OO" DESIGN



Why the quotation marks?

- C++ offers many ways to address the sort of problems covered by OO design
- Not all of these are considered "OO" techniques
- Let's look at an example



TEMPLATES AND OO CAN SOLVE SIMILAR PROBLEMS



**Generic Programming
Should “Just” Be
Normal Programming**

— Bjarne Stroustrup

Using templates instead of inheritance and virtuals



- OO and templates can be used for many of the same problems
- In the spirit of Bjarne's quote, new C++ features, like C++20 Concepts, were intentionally designed to be familiar for replacing object-oriented code
- Let's review an example we looked at before

OO and Templates can solve similar problems



- Consider the following OO code

```
struct Animal {  
    virtual string name() = 0;  
    virtual string eats() = 0;  
  
};  
  
class Cat : public Animal {  
    string name() override { return "cat"; }  
    string eats() override { return "delicious mice"; }  
  
};  
// More animals...  
  
int main() {  
    unique_ptr<Animal> a = make_unique<Cat>();  
    cout << "A " << a->name() << " eats " << a->eats();  
}
```



Do we need to use OO?

- Not really

```
struct Cat {  
    string eats() { return "delicious mice"; }  
    string name() { return "cat"; }  
};  
// More animals...  
  
int main() {  
    auto a = Cat();  
    cout << "A " << a.name() << " eats " << a.eats();  
}
```


That was a lot simpler but...



- We lost the understanding that a is an animal
- a could have the type House or int and we might not find out that something went wrong until much later when we did something that depends on a being an animal
- What we need is a way to codify our expectations for a without all of the overhead and complexity of creating a base class



Concepts

- Concepts play the analogous role for generic programming that base classes do in object oriented programming
- A concept explains what operations a type supports
- The following concept encapsulates the same info as the base class

```
template<typename T>
concept Animal = requires(T a) {
    { a.eats() } -> convertible_to<string>;
    { a.name() } -> convertible_to<string>;
};
```

Now, we can ensure that a represents an animal



- With the above concept defined, we can specify that a must satisfy the Animal concept, and the compiler will not let us initialize it with a non-Animal type like House or int

```
int main() {  
    Animal auto a = Cat();  
    cout << "A " << a.name() << " eats " << a.eats();  
}
```



Let's compare

- <https://godbolt.org/z/cWc6aM>
- As you can see, the definition of Cat and the client code in main() look very similar in both
- This follows a principle enunciated by Bjarne Stroustrup
 - “Generic Programming should just be Normal Programming”

Usage is almost identical to before



```
Animal auto a = Cat();  
cout << "A " << a.name()  
      << " eats " << a.eats();
```



How does this compare?

- Performance is better
 - Objects created on stack
 - No virtual dispatch
- No inheritance
 - Makes it easier to adapt classes to our code without risking “spaghetti inheritance”
 - On the other hand, it weakens type safety
 - Pacman is not an animal but eats and has a name
- No runtime polymorphism
 - The following is legal if `Animal` is a class but not if it is a concept (Why?)
 - `set<unique_ptr<Animal>> zoo;`



A real world example

- Suppose C++ didn't have mutexes
 - It didn't until C++11
- How would we design them?
- Let's look at how Java does it
- Java uses inheritance and virtual methods



Java-style mutexes

```
struct lockable {  
    virtual void lock() = 0;  
    virtual void unlock() = 0;  
};  
struct mutex: public lockable {  
    void lock() override;  
    void unlock() override;  
};  
struct lock_guard {  
    lock_guard(lockable &m) : m(m) { m.lock(); }  
    ~lock_guard() { m.unlock(); }  
    lockable &m;  
};
```




Using our Java-style mutex

```
mutex m;  
  
void f() {  
    lock_guard lk(m);  
    // do stuff  
}
```

Wait, that's not how C++ mutexes work!



- C++ mutexes do not inherit from a lockable base class
- The C++ committee decided to use templates instead of the virtual override approach taken by Java
- Since mutexes are frequently used in performance-critical code, this was undoubtedly the right choice
- Let's take a look



C++-style mutexes

```
struct mutex {  
    void lock();  
    void unlock();  
};
```

```
template<typename T>  
struct lock_guard {  
    lock_guard(T &m) : m(m) { m.lock(); }  
    ~lock_guard() { m.unlock(); }  
    T &m;  
};
```

Using our C++-style mutex is typically unchanged



```
// Exactly the same as before!
```

```
mutex m;
```

```
void f() {  
    lock_guard lk(m);  
    // do stuff  
}
```

Class Template Argument Deduction



- Note the following line depended on C++17 CTAD
 - `lock_guard lk(m);`
- CTAD infers the template arguments for `lock_guard<mutex>` from the constructor similarly to how Function Template Argument Deduction infers the template arguments for function templates
- CTAD can often be useful in this way when using templates instead of virtuals. E.g., if `tp1` and `tp2` are of type `time_point<C, duration<R>>`
 - `duration d = tp1 - tp2; // duration<R>`



So many choices :/

- As we saw above, there is usually a choice between base classes/virtuals and templates/concepts
- As we will see below, this is only the beginning
- C++ supports many approaches for "object-orientation"
- How can we make sense of this?
- We will need an understanding of OO design best practices



THE SOLID PRINCIPLES (H/T TONY VAN EERD)

Popularized by "Uncle Bob"

Robert C Martin ~2004



- **Single Responsibility Principle**
- **Open/Closed Principle**
- **Liskov Substitution Principle**
- **Interface Segregation Principle**
- **Dependency Inversion Principle**



Single Responsibility Principle

- How often do you see a class that has members organized into subsets?

```
class Camera
{
    CameraId id;
    DevicePath path;

    int bitdepth;
    Resolution resolution;
    int gain;
    int exposure; // units?

    Pose pose;
    Calibration * calibration;



    int binarizationThreshold;
    int sharpness; // UI only
    int multicamEdgeThreshold;

    Image baseImage;
    Image negativeImage;
    Image maskToUse;
    Image reverseMask;

    Device * device;
    INativeCamera * camera;

    ImagePoller * poller;
};
```

These are separate pieces pulled together by association



What if we wanted to use just one piece?



- Can I pose (position) things other than a camera?
- Does posing really have anything to do with resolution?
- What if a camera supported multiple resolutions but only a single pose?
- We can't do that with this class without breaking the Open-Closed principle
 - Which we'll learn about momentarily, but in this context it means "without breaking existing code"



Better: Give Image its own class



```
struct Image
{
    unsigned char * pixels;
    Format format; // RGB vs RGBA vs Gray vs ...
    int width;
    int height;

    Image();
    Image(int width, int height, Format format);
};
```

It's not just classes that mix concerns



- What about functions?
- Many functions are thousands of lines long
 - E.g., <https://github.com/llvm/llvm-project/blob/b07aab8fc1088ef66ecbe2befc3ef7e3936a390e/clang/lib/Parse/ParseExprCXX.cpp#L154>
- ```
void myFunction(Quux q) {
 // Locals
 Pass p(1);
 Bar b(p);
 Blah l(q);
 // Setup the frabulator
 Frabulator f;
 f.x();
 f.y = "foo";
 // Loop through the sneetches
 for(auto s: f.sneetches)
 ...
```

# How can I break a function into single-concern pieces?



- Breaking into several functions is typical
- But is sometimes difficult to share state
- ```
void myFunction(Quux q) {  
    // Locals  
    Pass p(1);  
    Bar b(p);  
    Blah l(q);  
    Frabulator f = setupFrabulator(b, p, l, q);  
    loopThroughSneetches(f, b, p, l, q);  
    ...  
}
```

Option 1: More, smaller functions



- Breaking into several functions is typical
- But is sometimes difficult to share state
- ```
void myFunction(Quux q) {
 // Locals
 Pass p(1);
 Bar b(p);
 Blah l(q);
 Frabulator f = setupFrabulator(b, p, l, q);
 loopThroughSneetches(f, b, p, l, q);
 ...
}
```





# Option 2: Use a functor

- Functions can be restructured as functors for better organization
- ```
struct myFunctionHelper {  
    myFunctionHelper(Quux q) : q(q), l(q) {}  
    Quux q;  
    Pass p(1);  
    Bar b(p);  
    Blah l;  
    // Methods now have access to variables  
    Frabulator setupFrabulator() { ... }  
    void loopThroughSneetches() { ... }  
    int operator() {  
        Frabulator f = setupFrabulator();  
        loopThroughSneetches(f);  
        ...  
    };  
    int myFunction(Quux q) { return myFunctionHelper(q)(); }  
};
```
- Good choice for really complex functions
- Supports powerful features like virtual functions and multiple entry points
- but high-friction for simpler cases

Option 3: Use lambda with capture lists



- While C++ doesn't have local functions *per se*, you can get the same organizing effect with local lambdas
- ```
int myFunction(Quux q) {
 Pass p(1);
 Bar b(p);
 Blah l;
 auto setupFrabulator = [&] { ... }
 auto loopThroughSneetches = [&](Frabulator f) { ... }
 Frabulator f = setupFrabulator();
 loopThroughSneetches(f);

 ...
};
```
- Good for "goldilocks" cases where a class is overkill but multiple responsibilities risk a "spaghetti" function



# Open/Closed Principle

- Software constructs should be open for extension but closed for modification
- Open for extension
  - Allows the addition of new capabilities over time
- Closed for modification
  - Don't break existing client code



# Inheritance and virtuals

- Of course, inheritance and virtual functions are a great way to extend classes
  - As are the template equivalents we discussed above
- Is it enough?
- Not quite
- It does not cover extremely common use cases for extensions

# The Problem: Client-side extension



- Suppose you are using a class hierarchy, and you wish the classes had a virtual method specific to the needs of your application
- Unfortunately, it probably doesn't because the class designer doesn't understand your application
- You may not be able to add them
  - Maybe they're not your classes
  - Maybe the virtuals you want only apply to your particular program, and it breaks encapsulation to clutter up a general interface with the particulars of every app that uses them



# The Visitor Pattern

- The Visitor Pattern is a way to make your class hierarchies extensible
- Suppose, as a user of the `Animal` class, I wished that it had a `lifespan()` method, but the class designer did not provide one
- We will fix that with the visitor pattern

# Class Creator: Make Your Classes Extensible



- Create a visitor class that can be overridden
- ```
struct AnimalVisitor {  
    virtual void visit(Cat &) const = 0;  
    virtual void visit(Dog &) const = 0;  
};
```
- Add an “accept” method to each class in the hierarchy
- ```
struct Animal {
 virtual void accept(AnimalVisitor const &v) = 0;
};
struct Cat : public Animal {
 virtual void accept(AnimalVisitor const &v)
 { v.visit(*this); }
 /* ... */
};
```



# Class User: Create a visitor

- Now, I create a visitor that implements the methods I wish were there
- ```
struct LifeSpanVisitor
    : public AnimalVisitor {
    LifeSpanVisitor(int &i) : i(i) {}
    void visit(Dog &) const { i = 10; }
    void visit(Cat &) const { i = 12; }
    int &i;
};
```




Using the visitor

- Now, I can get the lifespan of the Animal a I created above
- ```
int years;
a->accept(LifeSpanVisitor(years));
cout << "lives " << years;
```
- <https://godbolt.org/z/WbGe4z>



# Best practice

- If you are designing a class hierarchy where the best interface is unclear, add an `accept()` method as a customization point



# USING DUCK-TYPED VARIANTS FOR PERFORMANT, EXTENSIBLE OO

# Duck Typing





# Duck Typing

- OK. Not that
  - When I gave a talk on this in Shenzhen, I was worried what the translator would do with the last slide :/
- There is a saying

If it walks like a duck and quacks like  
a duck, then it is a duck

- Let's see if we can apply this to types



# Inheritance models “isA”

- As we’ve mentioned, inheritance is a model of the “isA” concept
- Duck typing gives a different notion of “isA”
- If a class has a walk() method and a quack() method, let’s not worry about inheritance and call it a duck

If it walks like a duck and quacks like a duck, then it isA duck



# Templates use duck typing

- `T square(auto x) { return x*x; }`
- `square(5); // OK`
- We don't require that the type of `x` inherits from a `HasMultiplication` class, simply that `operator*` makes sense to use here



# Concepts

- C++20 Concepts improve duck typing
- We could have a HasMultiplication concept that would do the trick without requiring any complex inheritance hierarchies





# Dynamic dispatch

- While C++ templates have always been duck typed, templates are used for compile-time dispatch
- By contrast, OO is used for dynamic dispatch
- Because duck typing is flexible and forgiving while remaining statically typesafe, people have asked whether we could use dynamically-dispatched duck typing as an alternative to inheritance



# Variants

- C++ has a lightweight "variant" abstraction that generalizes C unions
- A `variant<A, B>` can hold an A or a B but not both
- These variants will be the basis of an approach to OO that will
  - Often have much better performance than virtuals
  - More dynamic than our `Animal` concept (we can have a zoo with runtime dispatch)
  - Everything is value-based, no need to worry about references, `unique_ptr`, `RAII`, ...
  - Great support for Open-Closed extensibility



# Variants: Basic use

- But first, what is a variant?
- A variant is a lot like a tuple, but instead of holding all of its fields at once, it only contains one of them at a time
- Supports a very similar interface to tuples
- ```
variant<int, double> v = 3; // Holds int  
get<0>(v); // Returns 3  
get<1>(v); // Throws std::bad_variant_access  
v = 3.5;    // Now holds a double  
get<double>(v); // Returns 3.5
```
- You can also check what is in it

```
v.index; // returns 1  
holds_alternative<double>(v); // returns true  
holds_alternative<int>(v); // returns false
```

Using Duck-Typed Variants in place of OO



- Suppose we knew (at compile-time) all of the Animal classes
 - E.g., Cat and Dog
- However, we don't know the type of a particular animal until run-time
- Instead of inheriting from an abstract animal base class, we can have an animal variant
- `using Animal = variant<Cat, Dog>;`

How do I call a method on a variant?



- While `variant<Cat, Dog>` is a great way to store either a `Cat` or a `Dog`
- How can I simulate virtual functions and call the right `name()` method for the type it is holding?

The C++ standard library has a solution: `std::visit`



- **Warning:** Do not confuse with the Visitor Pattern we discussed earlier
- If `v` is a variant, and `c` is a callable, `visit(v, c)` calls `c` with whatever is stored in `v` as its argument
- Does this solve our problem of making variants behave like virtuals?
- Let's see

Dynamically calling our `Animal's name()` method



- `Animal a = Cat(); // a is a cat`
- `cout << visit(
 [](auto &x) { x.name(); },
 a);`
- Prints “Cat”, just like we want
- How does it compare to using virtuals or templates?

In some ways, it's the best of both world



- Almost as fast as templates
 - Since Animal is a single type that can hold a Cat or a Dog, we can just use animals by value instead of having to do memory allocations
- As dynamic as traditional OO
 - A `set<Animal>` works great
 - Unlike our Concepts version



Malleable (mallard?) typing

- Virtual functions need to exactly match what they override, so we couldn't, say, give Cat's eat() a defaulted argument
 - ```
struct Cat : public Animal {
 void eat(string prey = "mouse") override; // ill-formed
```
- With variants, that is not a problem
- As long as eats() is callable, we don't care about the rest
- ```
struct Cat {  
    void eat(string prey = "mouse");  
};  
visit([](auto &x) { x.eats(); }, a); // OK. Eats a mouse
```



Users can add methods

- Just like we discussed with the Visitor Pattern, users of the Animal type can add their own methods
- To do this, we will use the “overloaded pattern”



The overloaded pattern

- Define an overloaded class (you only need to do this once)
- ```
template<class... Ts>
struct overloaded : Ts... { using Ts::operator()...; };
```
- This inherits the function call operator of everything it is constructed with
- Let's make this clear by creating a `lifeSpan()` “method” like we did before
- ```
overloaded lifeSpan([](Cat &) { return 12 },  
                   [](Dog &) { return 10});  
// Get life span without knowing whether  
// a is a Cat or a Dog  
cout << visit(lifeSpan, a); // Prints 12
```
- Note that this idiom relies on CTAD and aggregates deducing the template arguments for you



Problems with Duck Typing

- The notation is much uglier than calling a virtual function
- While the flexibility is nice, duck typing reduces type safety
 - It cannot tell that a Shape's draw() method for drawing a picture is different than a Cowboy's draw() method for drawing a gun
- Variants always uses as much space as the biggest type
- Whenever we create a new kind of Animal, we have to add it to the variant, which can create maintenance problems



Best Practice

- Because it is ugly and not well-known, only prefer variant-based polymorphism over virtuals or templates when there is a clear benefit
 - In practice, I use it a lot, but less than I do virtual functions or templates

Liskov Substitution Principle



Subtype Requirement:

Let $\Phi(x)$ be a property provable about objects x of type T .

Then $\Phi(y)$ should be true for objects y of type S

where S is a subtype of T



Inheritance models "isA"

- Inheritance is one way of modeling subtyping
- A Deer isA Animal
- One would expect that anything that is true about animals is true about deer



Concepts model is A

- If Animal is a concept instead of a base class, we have seen that the same is true
- One other benefit of concepts is that inheritance only inherits methods, but concepts can specify almost arbitrary $\Phi(x)$ properties
- Tradeoff: Efficiency vs dynamism
 - Generally how to choose the approach

Tony says



Liskov Substitution Principle

```
struct Line
{
    explicit Line(LineSegment);
    explicit Line(Ray);
};
```

```
struct Ray
{
    explicit Ray(Line);
    explicit Ray(LineSegment);
};
```

```
struct LineSegment
{
    explicit LineSegment(Line);
    explicit LineSegment(Ray);
};
```

LSP

Litmus for explicit vs implicit

```
int func(Ray r);

int main()
{
    Line line = ...;
    return func(Ray(line));
}
```

Interface Segregation Principle



- No code should be forced to depend on methods it doesn't use
- (Martin) Suppose we have a fat "Job" class that has a bunch of methods that are only relevant to print jobs and other methods that are only relevant to stapling jobs
- If the stapling code takes a Job, it will needlessly only work with Jobs that also know about printing



Handling with OO

- This is often given as a motivation for using abstract base classes (interfaces)
- The concrete Job class implements the PrintJob and StapleJob interfaces
- This can be taken so far, getting into spaghetti inheritance and excessive multiple inheritance complexity

Concepts also handle this nicely



- The stapling code can only require what it needs to staple without exploding the type hierarchy
- However, you could also go too far with this too as an incoherent set of functions that each make different requirements of each job that is passed in
- Both of these are good reminders that architecture is more art than science

Dependency Inversion Principle



- This is sometimes paraphrased as "All programming problems can be solved with an extra layer of indirection"
- 😊



Basic idea

- *“ the most flexible systems are those in which source code dependencies refer only to abstractions, not to concretions ”*



Example

- Suppose you have a thumbnail service class that looks for pictures in S3 folders

```
class ThumbnailService {  
    S3Folder inputFolder;  
  
    ...  
};
```
- It is now coupled with the concrete S3 service instead of an abstract idea of a storage service, which is probably sufficient for this use
- Again, the indirection can be introduced either through inheritance/virtuals or template/concepts
- Usual performance/dynamism tradeoff

Solving with inheritance and virtuals



- ```
Class S3Folder : public Folder { ... };
class ThumbnailService {
 Folder &inputFolder;

 ...
};
```



# Solving with templates and concepts



- ```
template<Folder F> // Folder is a concept
class ThumbnailService {
    Folder &inputFolder;
    ...
};
ThumbnailService<S3Folder> ts;
ThumbnailService ts2(myS3Folder); // CTAD
```



HW 15-1

- *Static polymorphism* is another common idiom for using templates to achieve "OO-like" behavior at compile-time
- Read the static polymorphism section of <https://iamsorush.com/posts/static-polymorphism-cpp/>
- Reimplement our "animal example" from <https://godbolt.org/z/cWc6aM> using static polymorphism
- What is your opinion of this approach? Do the SOLID principles shed any light on that?



HW 15-2

- Convert our "train example" on Canvas to use the variant/visit approach instead of inheritance
 - E.g., ModelLocomotive will no longer inherit from Locomotive
 - **Note:** Just adapt the train_factory example, not the advanced_train_factory
- **Hint:**
 - First just get the trains working without the factory (6/10 pts)
 - Next, figure out how to make factory templates that work with variant/visitor (4/10)



HW 15-3: Extra Credit

- Choose some C++ program (e.g., from GitHub)
- Analyze it according to the SOLID principles
- What does it do well?
- What does it do badly?
- Can you suggest improvements?